

# Efficient Distributed Stochastic Dual Coordinate Ascent

---

Jeff Hajewski

Mingrui Liu

May 3, 2017

University of Iowa

# Problem Overview

---

## Problem Overview - The Primal Problem

Many machine learning problems can be formulated as the Regularized Finite Sum Minimization (RFSM) problem.

$$\min_{w \in \mathbb{R}^d} P(w) \tag{1}$$

where

$$P(w) = \frac{1}{n} \sum_{i=1}^n \phi(w^\top x_i, y_i) + \lambda g(w)$$

$$w, x_i \in \mathbb{R}^d, \text{ for } i = 1, \dots, n$$

$$y_i \in \mathbb{R}, \text{ for } i = 1, \dots, n$$

$\phi(z, y)$  is convex in  $z$

$g(w)$  is convex in  $w$

# Problem Overview - The Dual Problem

Summary of the dual problem

The two key papers influencing our work are:

- Stochastic Dual Coordinate Ascent (SDCA) [Shalev-Shwartz and Zhang, 2013]
- Distributed SDCA [Yang, 2013, Yang et al., 2013]

# What is SDCA?

- **SDCA** - randomly pick a coordinate axis, find update that best improves the objective
- **Distributed SDCA** - randomly pick  $k$  coordinate axes, simultaneously find updates that best improve the objective (independently)

Two approaches:

- Naive approach: simply parallelize all tensor operations

Two approaches:

- Naive approach: simply parallelize all tensor operations
- Better: mimic the distributed approach by [Yang, 2013] on a GPU



# Implementation

---

Two key areas of concern:

- Memory
  - Allocation
  - Communication

Two key areas of concern:

- Memory
  - Allocation
  - Communication
- Code Abstraction (wrappers)

# Dealing with Memory Allocation

Naive approach:

```
void MemSync::PushToGpu(const vector &x) {  
    double *dx;  
    int n_bytes = sizeof(double) * x.size();  
  
    // Allocate GPU memory  
    cudaMalloc((double**) &dx, n_bytes);  
  
    // Copy data to GPU  
    cudaMemcpy(dx, &x[0], n_bytes, cudaMemcpyHostToDevice)  
}
```

On a small dataset(200 points in  $\mathbb{R}^3$ ) we hit over 200k allocations, which comprised nearly **95%** of the GPU compute time ( $\approx 13$  seconds).

**Can we do better?**

**Can we do better?**

**Yes!**

# Dealing with Memory Allocation

Better approach:

```
void MemSync::PushToGpu(const vector &x) {  
    int n_bytes = sizeof(double) * x.size();  
  
    // Copy data to GPU  
    cudaMemcpy(MemSync::dx_, &x[0],  
               n_bytes, cudaMemcpyHostToDevice);  
}
```

The use of static class pointers reduced the 200k allocations down to only **3 memory allocations** which comprised only **0.03%** compute time ( $\approx 180\mu s$ ).

**What about the cost of communication?**



Copying data to and from the GPU is the next most expensive operation

- About 50% of the compute time, or 768 ms (using the same toy dataset)
- Mostly unnecessary!

# Communication Costs

Consider the following algorithm using the GPU:

$$1: \Delta\omega \leftarrow f(\mathbf{x}, \omega)$$

$$2: \omega \leftarrow \omega + \Delta\omega$$

To handle this efficiently we should:

- Reuse  $\omega$  in step 2 since we already moved it to the GPU for step 1
- Perform step 2 on the GPU since the data is already there. No need to pull it off and then move it back to the GPU

The sad reality is this is quite complicated.

- Lots of book-keeping
- Are there edge cases?
- Need to watch out for memory leaks. Remember, **no** GC!

**What's the solution?**

**What's the solution?**  
**Wrappers!**

# Wrappers

We use wrappers (also known as decorators) to add additional functionality to our code. For example,

```
double VectorDotProduct(const vector &x,  
                        const vector &y) {  
    #ifdef GPU  
        return VectorDotProduct_gpu(x, y);  
    #else  
        return VectorDotProduct_cpu(x, y);  
    #endif  
}
```

To handle the flow of data from GPU to CPU, as well as book-keeping, we can use something like this:

```
class Data {  
    enum DataLocation { Gpu, Cpu };  
  
    // Pointer to GPU memory  
    double *dx_  
    // Local reference (RAM)  
    Eigen::VectorXd x_  
    // DataLocation::Gpu or DataLocation::Cpu  
    DataLocation location_  
};
```

**What about results?**



We are still finalizing results. There is a lot of non-trivial structural code behind this.

- Loading data
- Algorithm
- Algorithm performance tracking



Bottou, L. (2010).

**Large-scale machine learning with stochastic gradient descent.**

In *Proceedings of COMPSTAT'2010*, pages 177–186.  
Springer.



Boyd, S., Parikh, N., Chu, E., Peleato, B., and Eckstein, J. (2011).

**Distributed optimization and statistical learning via the alternating direction method of multipliers.**

*Foundations and Trends® in Machine Learning*,  
3(1):1–122.



Johnson, R. and Zhang, T. (2013).

**Accelerating stochastic gradient descent using predictive variance reduction.**

In *Advances in Neural Information Processing Systems*, pages 315–323.



Lian, X., Huang, Y., Li, Y., and Liu, J. (2015).

**Asynchronous parallel stochastic gradient for nonconvex optimization.**

In *Advances in Neural Information Processing Systems*, pages 2737–2745.



Nemirovski, A., Juditsky, A., Lan, G., and Shapiro, A. (2009).

**Robust stochastic approximation approach to stochastic programming.**

*SIAM Journal on optimization*, 19(4):1574–1609.



Shalev-Shwartz, S. and Zhang, T. (2013).

## **Stochastic dual coordinate ascent methods for regularized loss minimization.**

*Journal of Machine Learning Research*, 14(Feb):567–599.



Shalev-Shwartz, S. and Zhang, T. (2014).

### **Accelerated proximal stochastic dual coordinate ascent for regularized loss minimization.**

In *ICML*, pages 64–72.



Xiao, L. and Zhang, T. (2014).

### **A proximal stochastic gradient method with progressive variance reduction.**

*SIAM Journal on Optimization*, 24(4):2057–2075.



Yang, T. (2013).

### **Trading computation for communication: Distributed stochastic dual coordinate ascent.**

In *Advances in Neural Information Processing Systems*, pages 629–637.



Yang, T., Zhu, S., Jin, R., and Lin, Y. (2013).

**Analysis of distributed stochastic dual coordinate ascent.**

*arXiv preprint arXiv:1312.1031.*